

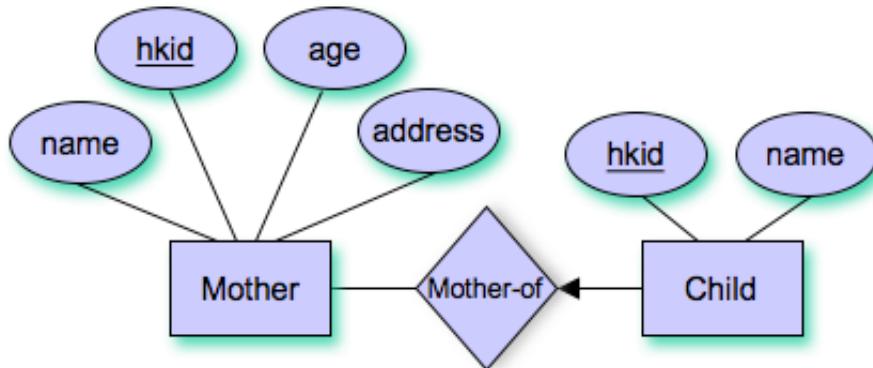
Entity-Relationship Model

- A key or candidate key is a minimal set of attributes to uniquely identify an entity

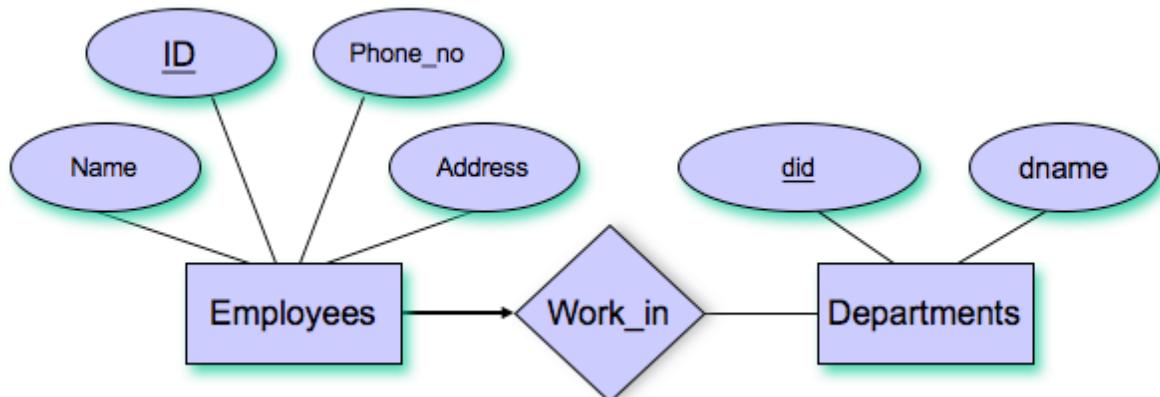
Key Constraints

One-to-many

- One-to-many constraint from A to B
- An entity in B can be associated with at most one entity in A

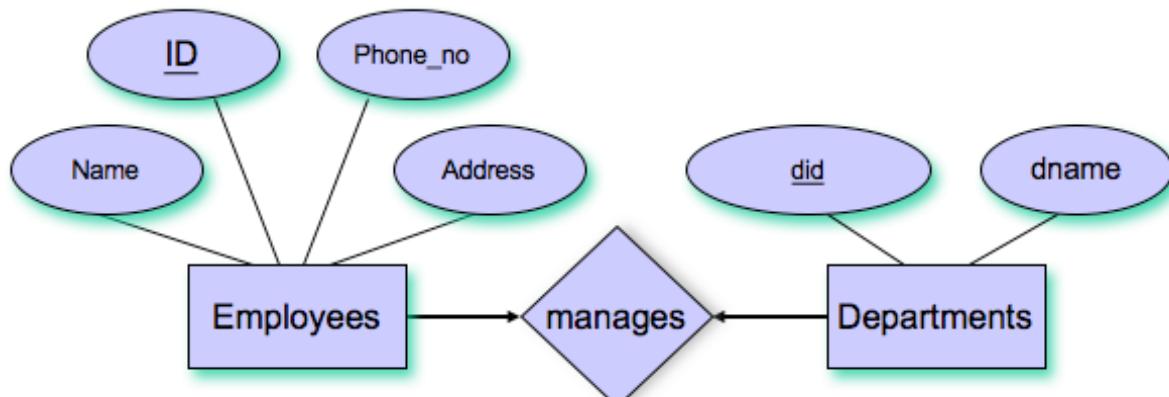


Many-to-one



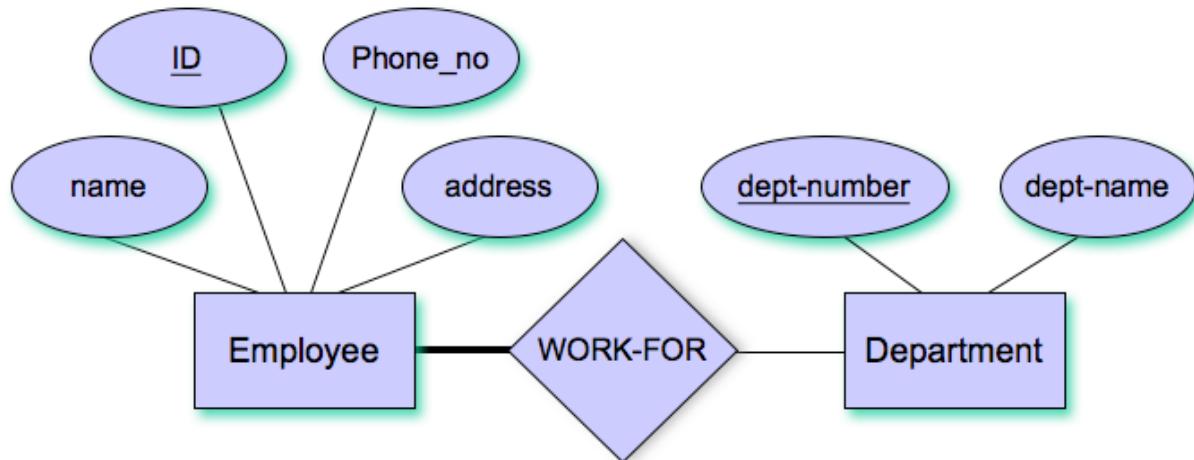
One-to-one

- An entity in A/B is related to at most one entity in B/A



Participation Constraints

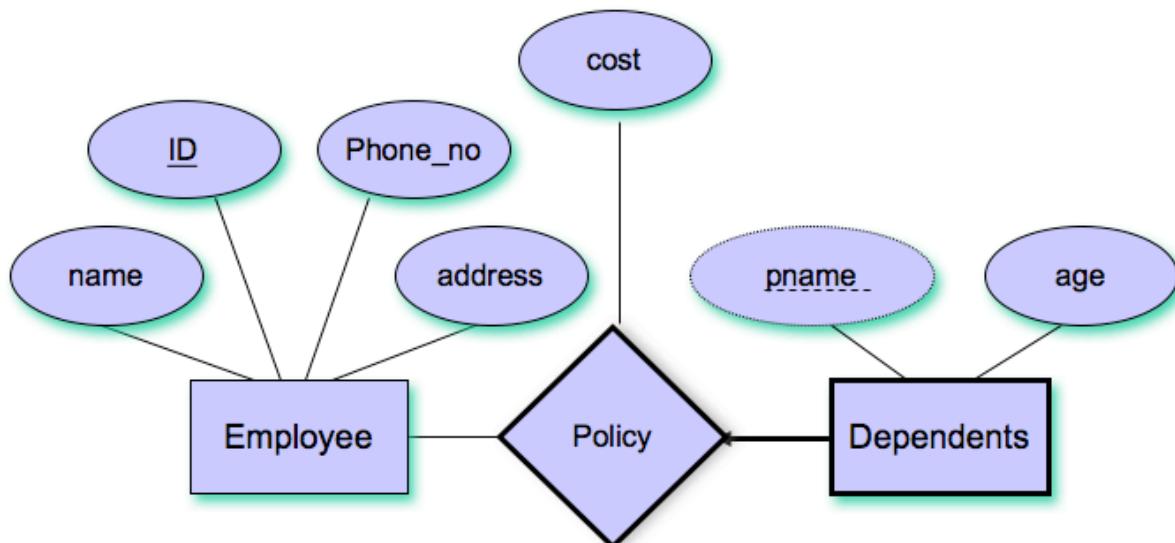
- Total: each entity in entity set must participate in **at least one** relationship
- Partial: an entity in entity set may not participate in a relationship



- Example of total participation

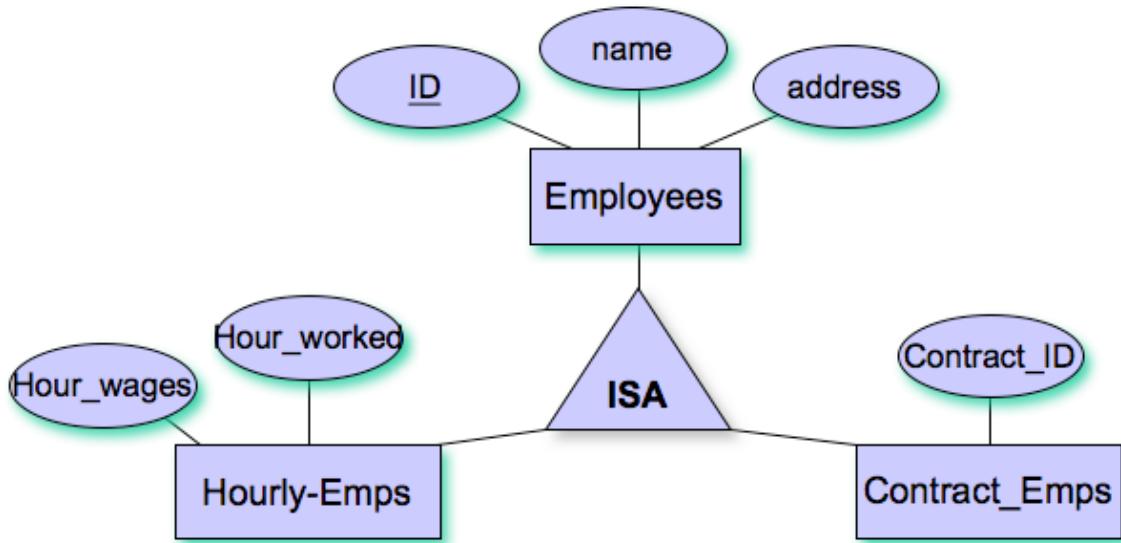
Weak Entities

- Without super key
- It may not be distinguishable from others **without associations** with entities in other entity sets



Class Hierarchies

- Classify entities in an entity set into sub-classes



Relational Algebra

Basic Operation

Select σ

Project Π

Union \cup

Set different –

Product \times

Rename ρ

Division $/$

Example

- Select Relation $a = b$ and $d > 5 = \sigma_{a=b \wedge d>5}(R)$
- $\Pi_{f_name, l_name}(\sigma_{dno=4 \wedge salary>25000}(Employee))$
- Rename $E.sid$ to $C = \rho(C(sid \rightarrow \text{identity}), E)$
- Condition Join = $R \bowtie_c S = \sigma_c(R \times S)$

SQL

Basic SQL Query

```

1  SELECT [DISTINCT] target-list
2  FROM relation-list
3  WHERE qualification

```

- Equal Join with tables

```

1 SELECT S.sname
2 FROM Sailors S, Reserves R, Boat B
3 WHERE S.sid = R.sid AND R.bid = B.bid

```

- Regex in SQL

```

1 SELECT S.age, age1=S.age-5, 2*S.age AS age2
2 FROM Sailors S
3 WHERE S.sname LIKE 'B_%B'

```

`LIKE` for string matching, " _ " means any characters, "% " means arbitrary

- Set-manipulation construction

- `UNION` \cup
- `INTERSECTION` \cap
- `EXCEPT` —

```

1 SELECT S.sid
2 FROM Sailors S, Boats B, Reserves R
3 WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
4 INTERSECT
5 SELECT S.sid
6 FROM Sailors S, Boats B, Reserves R
7 WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='green'

```

```

1 SELECT R.sid
2 FROM Boats B, Reserves R
3 WHERE R.bid=B.bid AND B.color='red'
4 EXCEPT
5 SELECT R.sid
6 FROM Boats B, Reserves R
7 WHERE R.bid=B.bid AND B.color='green'

```

- Correlated Nested Query

```

1 SELECT S.sname
2 FROM Sailors S
3 WHERE EXISTS(SELECT * FROM Reserves R WHERE R.bid=103 AND S.sid=R.sid)

```

- `EXISTS` to test for nonempty
- `IN` operator specified **multiple values** in `WHERE` clause

```

1 SELECT B.bname
2 FROM Boats B
3 WHERE B.color IN ('red', 'blue', 'green')

```

- **Set-comparison** operators
- `ops ANY` or `ops ALL`

```

1 SELECT *
2 FROM Sailors S
3 WHERE S.rating > ANY (SELECT S2.rating FROM Sailors S2 WHERE
                           S2.sname='Horatio')

```

- Division in SQL

```

1 SELECT S.sname
2 FROM Sailors S
3 WHERE NOT EXISTS
4     ((SELECT B.bid
5         FROM Boats B)
6         EXCEPT
7         (SELECT R.bid
8             FROM Reserves R
9             WHERE R.sid=S.sid))

```

SQL Aggregate Operators	Description
<code>COUNT(A)</code>	Number of values in A column
<code>SUM(A)</code>	Sum of all values in A column
<code>AVG(A)</code>	Average of all values on A column
<code>MAX(A)</code>	Maximum value in the A column
<code>MIN(A)</code>	Minimum value in the A column

- Correct way to use the above operators

```

1 SELECT S.name
2 FROM Sailors S
3 WHERE S.age > (SELECT MAX (S2.age)
4                         FROM Sailors S2
5                         WHERE S2.rating = 10)

```

- `GROUP BY` similar as for-loop

```

1 SELECT S.rating, MIN(S.age)
2 FROM Sailors S
3 GROUP BY S.rating
4
5 For i = 1, 2, ..., 10:
6     SELECT MIN(S.age)
7     FROM Sailors S
8     WHERE S.rating = i

```

- Columns appeared in `GROUP BY` should use `HAVING`
- `CREATE VIEW` for creating virtual table based on result SQL statement

```

1 CREATE VIEW Temp AS
2     SELECT S.rating, AVG (S.age) AS avgage
3     FROM Sailors S
4     GROUP BY S.rating
5
6 DROP VIEW temp

```

Schema Refinement

Functional Dependencies

```

1 Let X and Y be nonempty sets of attributes in R
2 An instance r of R satisfies the FD X->Y if
3 If t1.X = t2.X, then t1.Y = t2.Y

```

Trivially Preserved

```

1 If any two rows never have the same value for a in a->b
2 Then a->b is trivially preserved

```

Trivially Dependency

```

1 If right hand side of arrow is subset of that on left hand side a->b
2 Then a->b is a trivial dependency

```

Closure of set

- Given a set F, the set of all FDs implied is called the closure of F, denoted as F^+

Armstrong's Axioms and additional rules

- **Reflexivity:** if $Y \subseteq X$, then $X \rightarrow Y$

Augmentation: if $X \rightarrow Y$, then $XZ \rightarrow YZ$

Trasitivity: if $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$

Union: if $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$

Decomposition: if $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$

Boyce-Codd Normal Form (BCNF)

R – a relation schema

F – set of functional dependencies on R

R is in BCNF if for any $X \rightarrow A$ in F ,

- $X \rightarrow A$ is a trivial functional dependency, i. e. $A \subseteq X$
OR
- X is a superkey for R

Third Normal Form

- If R is in BCNF, then it is also in 3NF since $3NF \subset BCNF$

R – a relation schema

F – set of functional dependencies on R

R is in 3NF if for any $X \rightarrow A$ in F ,

- $X \rightarrow A$ is a trivial functional dependency, i. e. $A \subseteq X$
OR
- X is a superkey for R
OR
- A is part of some key for R

Decomposition

1. Remove the rule $X \rightarrow A$ that violates the condition from relation schema R
2. Create a new realtional schema XA

Lossless Join Decomposition

R – a relation schema

F – set of functional dependencies on R

The decomposition of R into relations with attribute sets R_1 , R_2 is lossless – join iff

$$(R_1 \cap R_2) \rightarrow R_1 \in F^+$$

OR

$$(R_1 \cap R_2) \rightarrow R_2 \in F^+$$

$R_1 \cap R_2$ is a superkey for R_1 or R_2

Dependency preserved

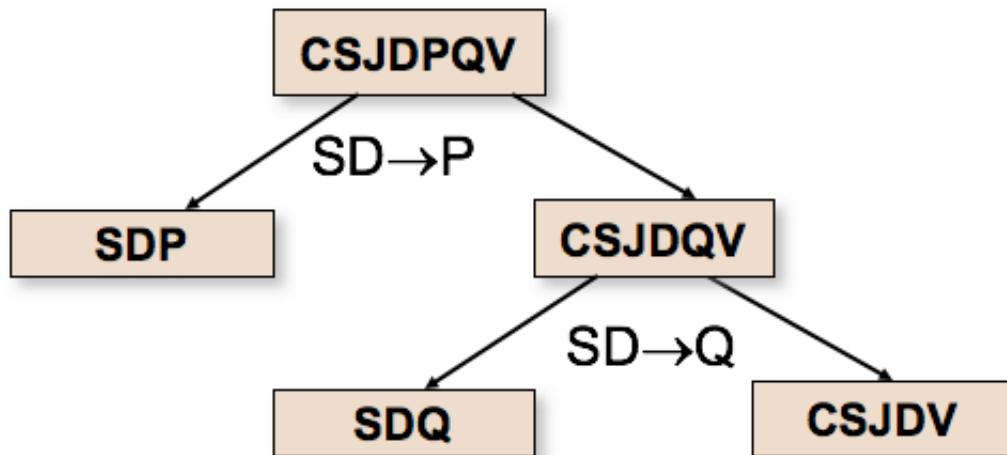
$$(F_1 \cup F_2)^+ = F^+$$

- Possible to obtain lossless-join decomposition into collection of BCNF relation schemas
 $\text{non-BCNF} \rightarrow \text{BCNF}$ but **NOT** guaranteed dependency-preserving
- Always exists a dependency-preserving, lossless-join decomposition into collection of 3NF relation schemas
 $\text{non-3NF} \rightarrow \text{3NF}$

BCNF Decomposition

Suppose $X \rightarrow A$ is a FD that violates the BCNF condition

1. Decompose R into XA and $R - A$
2. Repeat until all relations become BCNF

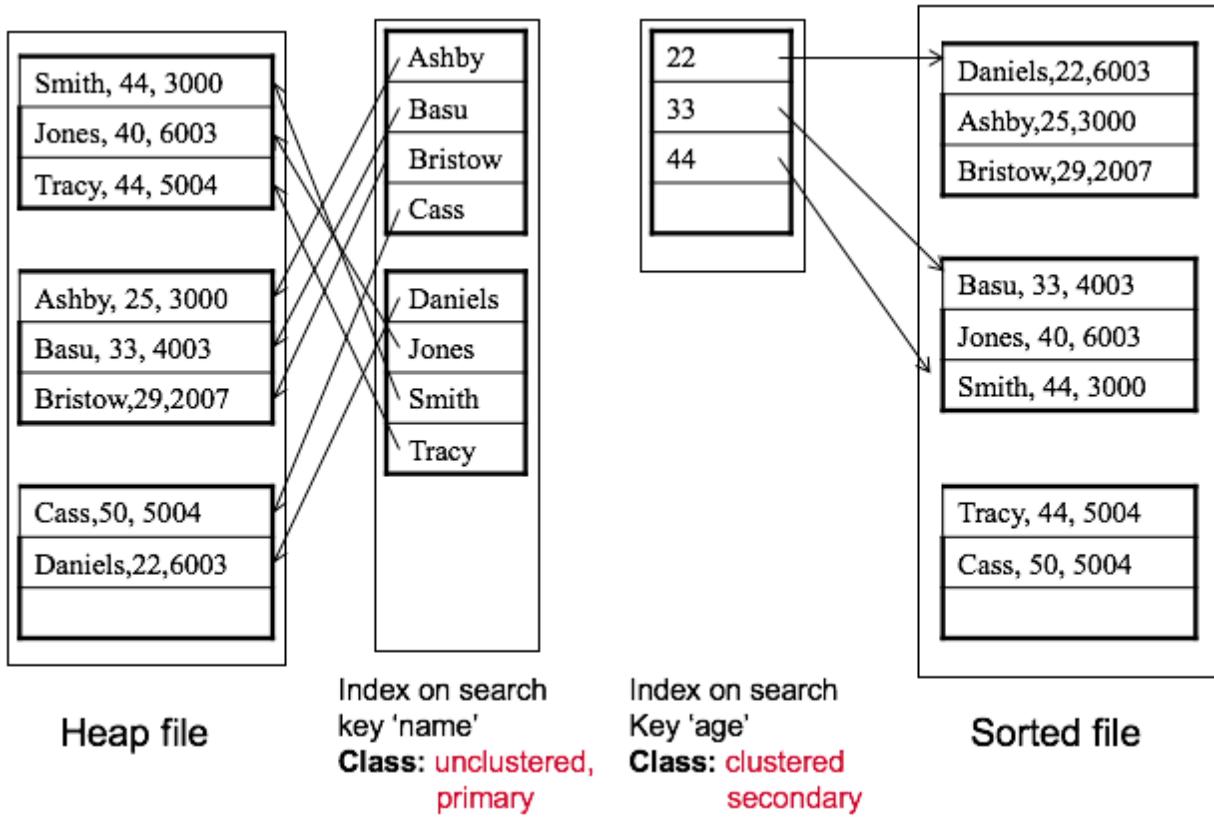


Canonical Cover

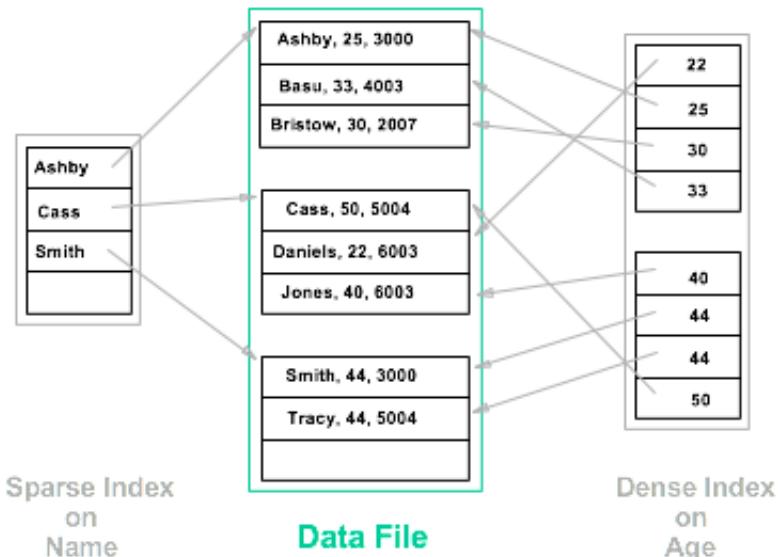
- A **minimal and equivalent** set of functional dependency

Storage and Index

- Index on file speeds up selections on **search key fields**
- Search key can be any subset of fields of relation



Primary index	Contains primary key in search key
Secondary index	Does not contain primary key in search key
Clustered index	Order of data records close to order of data entries
Unclustered index	
Dense index	At least one data entry per search key value
Sparse index	Every sparse index is clustered



- Unclustered must be also in **dense**

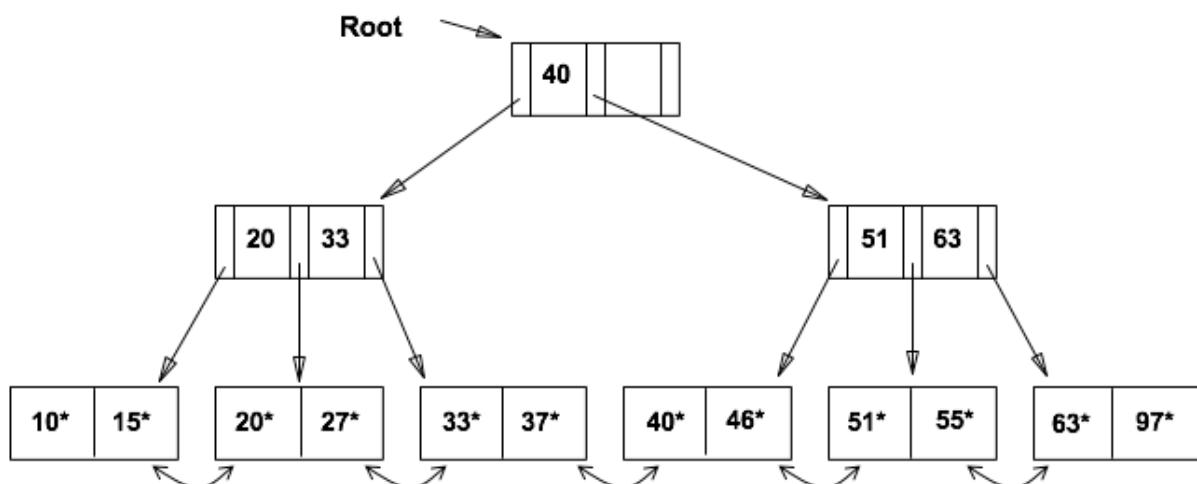
- Primary: each data entry k^* points to **single record** that contains k
- Secondary: each data entry k^* points to **all records** that contains k
- **Clustered** must be also in **sparse**
 - Sort both data file and index file on search key
 - Each data entry k^* points to the **first record** that contains k
 - Overflow pages may be needed for inserting, so the order is **closed to sorted**

Composite Search Keys

Equality query	Every field value is equal to a constant value
Range query	Some field value is not a constant

Tree-Structured Indexing

- Example of B+ tree with **order 1**



For root node, we require $1 \leq n \leq 2d$, where d is the order

For non – root node, we require $d \leq n \leq 2d$, where d is the order

Cost for searching in B+ tree

- Let h be the height of B+ tree, then we have to access $h + 1$ pages to reach leaf node
- Let f be the average number of pointers in node (fanout for internal node)
 - Level 1 with height 0 = 1 page = f^0 page
 - Level 3 with height 2 = $f \times f$ page = f^2 page
- Suppose there are d data entries, so there are $\frac{d}{(f-1)}$ leaf nodes and $h = \log_f(\frac{d}{f-1})$
- Example for calculation
 - Typical order = 100, Typical fill-factor = 67%
 - Average fanout $f = \frac{100}{67\%} = 133$
 - Given there are 10000000 data entries, $h = \log_{133}(\frac{10000000}{133-1}) < 4$

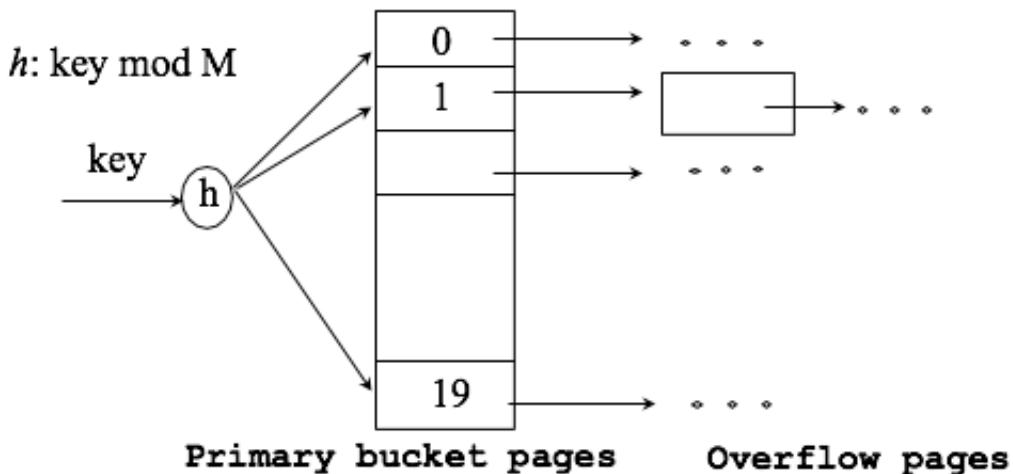
- Therefore, the cost is 5 pages read

Extensible Hashing

- Given a search key value k , we can find the bucket where data entry k^* is stored
- The value of hash function $h(k)$ is address of desired bucket

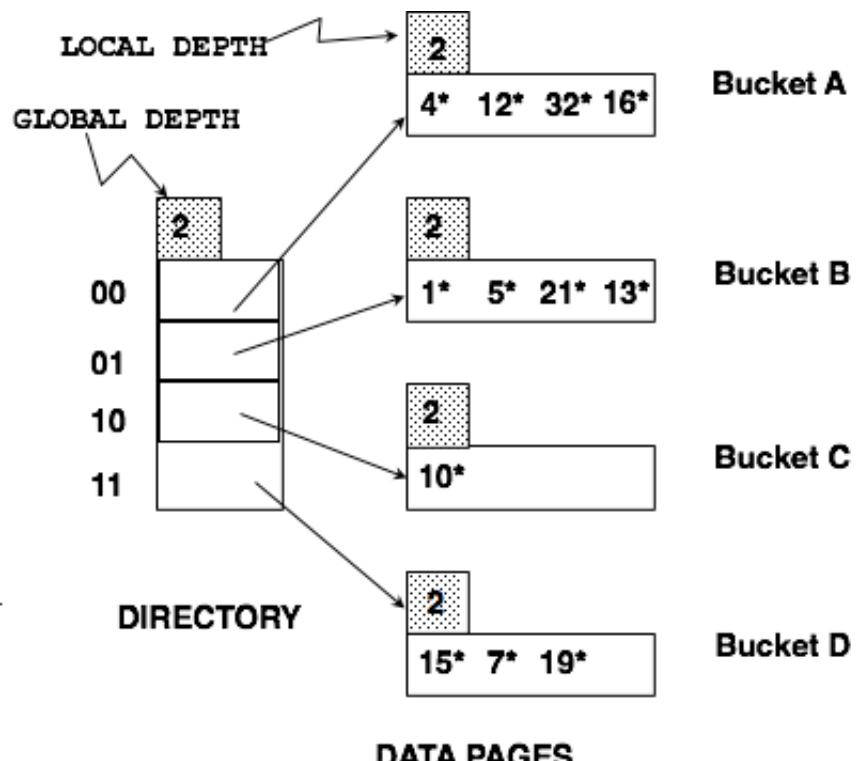
Hash-based indexes are the best for **equality selections** and they do **NOT** support range searches

Static Hashing



- Long overflow chains can be developed and degraded performance

Extensible Hashing



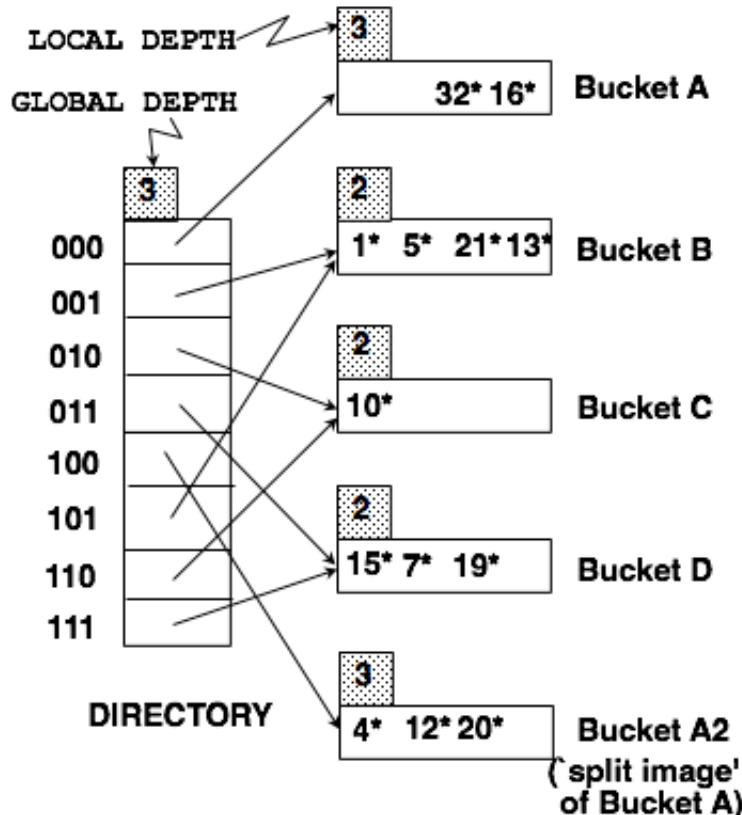
- To avoid re-hashing to re-organize file by doubling numbers of buckets, use directory of

- points to buckets
- Doubling here means increasing size of directories

Suppose we have a hash function $h(r)$ and directory is array of size 4

To find the bucket for r , take last x bits of $h(r)$ where x is number of global depth

- For inserting 20 in the above example, it will cause overflow and directory doubling is required



- Split Bucket A into 2 buckets and we compare 3rd bit from right in $h(r)$ to decide A or A_2
- Other buckets will remain unchange and 2 directories are pointed to bucket for *Global Depth > Local Depth*
- Least significant bits** are used in directory to allow for doubling via copying

Query Evaluation

- Buffer page** is in memory area and stored data page for processing
- Data file often needed to be **sorted** based on search key to speed up query evaluation

Suppose a file contains M pages and B buffers are available in memory

$$\text{File Sorting} = 2 \times M \times (\lceil \log_{B-1} \frac{M}{B} \rceil + 1) \text{ page accesses}$$

- Access Path** is method of retrieving tuples
- Selectivity** of an access = number of pages retrieved (index pages + data pages)
- Consider only I/O cost = **number of pages** that are read/written

Example on selection

Given the following schema :

Sailors(sid : integer, sname : string, rating : integer, age : real)

Reserves(sid : integer, bid : integer, day : dates, rname : string)

rname and sid may refer to different persons

Assuming the following sizes

Sailors : 500 pages, 80 tuples/page, 50 bytes/tuple

Reserves : 1000 pages, 100 tuples/page, 40 bytes/tuple

Consider selection : $\sigma_{R.\text{attribute}=\text{value}}(R)$ and M denotes as *Reserves* pages

No index and unsorted = M pages I/Os for sequential scanning

No index but sorted = $\log_2 M + M \times 10\%$ pages I/Os for binary searching + scanning

Above assumption = 10% tuples satisfy the condition

B + Tree index on R. attribute

1. Search tree to find first data entry that points to qualifying tuple of R
2. Scan leaf page to retrieve all data entries (**not needed for clustered index**)
3. Follow the pointers to corresponding tuple of R for each data entry

Cost for B + tree index

Clustered index	Unclustered index
1. Height of B+ tree, example: 3	1. Height of B+ tree, example: 3
2. Scan for numbers of leaf nodes	2. Scan data file: $1000 \text{ pages} \times 10\% = 100$
3. Costs: 10000 I/Os in worst case, each tuple in a page	

Hash index on R. attribute

1. 0 I/O for calculating hash value
2. 1 I/O if directory does not fit in memory, else 0 I/O
3. Retrieve bucket page pointed by directory entry, 1.2 I/Os for overflow pages
4. Depend on how many pages pointed and qualifying tuples

Example on projection

Assuming the following sizes

Sailors : 500 pages, 80 tuples/page, 50 bytes/tuple

Reserves : 1000 pages, 100 tuples/page, 40 bytes/tuple

```
1 | SELECT DISTINCT R.sid, R.bid  
2 | FROM Reserves R
```

- General method
 - Scan relation R and discard unwanted attributes
 - Eliminate duplicates ← Expensive operation
- Projection based on sorting
 1. Scan R and write sid and bid of each tuple to temporary file T
 2. Sort T based on attributes
 3. Scan sorted file and compare the adjacent tuples and discard duplicates
- Assume T has 250 pages and 20 buffers available
 1. $1000 + 250 = 1250$ I/Os
 2. $2 \times 250 \times (\lceil \log_{20-1} \frac{250}{20} \rceil + 1)$
$$File\ Sorting = 2 \times M \times (\lceil \log_{B-1} \frac{M}{B} \rceil + 1) \text{ page accesses}$$
 3. 250 I/Os

Example on join operation

Assuming the following sizes

Sailors : 500 pages, 80 tuples/page, 50 bytes/tuple

Reserves : 1000 pages, 100 tuples/page, 40 bytes/tuple

```
1 | SELECT *  
2 | FROM Reserves R, Sailors S  
3 | WHERE R.sid=S.sid
```

Block Nested Loop Join

- Assume there are enough buffers to hold $B + 2$ pages
1 page is reserved to read Sailors and 1 page is reserved for output
 - for each block P of Reserves*
 - for each page q of Sailors*
 - for each r ∈ P and s ∈ q such that r.sid = s.sid*
 - add < r, s > to the result*
- Assume $B = 100$, then Reserves R contains 10 blocks
 1. Scan Reserves R = 1000 I/Os
 2. Scan Sailors for each block = 500 I/Os
 3. Total number of I/Os = $1000 + 10 \times 500 = 6000$ I/Os

Possible to sway Reserves and Sailors relations for smaller cost

- If buffer size is large enough to store smaller relation + 2 pages, the cost will be $R + S = 1500$ I/Os

Index Nested Loop Join

- Assume an index on sid of Sailors relation

*for each r ∈ Reserves do
for each s ∈ Sailors where r.sid = s.sid
add < r, s > to the result*

- Cost for hash index

1. Scan Reserve: 1000 I/Os
2. For each tuple in Reserves, overflow 1.2 I/Os to get the matching bucket page
3. For each matching Sailors data entry, retrieve Sailors tuple for 1 I/O since sid is primary key
4. Each block of Reserves contains 100 tuples

- Total cost = $1000 + 100 \times 1000 \times (1 + 1.2) = 221000$ I/Os

- Assume an index on sid of Reserves relation

*for each s ∈ Sailors do
for each r ∈ Reserves where r.sid = s.sid
add < r, s > to the result*

- Cost for hash index

1. Scan Sailors: 500 I/Os
2. For each tuple in Sailors, overflow 1.2 I/Os to get the matching bucket page
3. For each matching Reserves data entry, retrieve Reserves tuples

Estimation on average = 100000 reservations for 40000 sailors = 2.5 reservations/sailor

- Consider **clustered index**, 2.5 reservations likely on same page

Total cost = $500 + 40000 \times 1.2 + 40000 \times 1 = 88500$ I/Os

- Consider **unclustered index**, 2.5 reservations not likely on same page

Total cost = $500 + 40000 \times 1.2 + 40000 \times 2.5 = 148500$ I/Os

Sort-Merge Join $R \bowtie_{i=j} S$

- Sort R, S on the join column, then scan them to perform merging on join column
- R is scanned once and each S group is scanned once per matching R tuple
- Total cost = cost for sorting two relations + joining two relations

Terminologies

- n_r : number of tuples in relation r

- f_r : number of tuples of r in one block
- $b_r = \lceil \frac{n_r}{f_r} \rceil$ if tuples are stored together
- S_r : size of tuple of relation r
- $V(A, r)$: number of distinct values that appear in relation r for attribute A
- $SC(A, r)$: average number of records that satisfy an equality condition on attribute A
 - $SC(A, r) = 1$ if A is a key of r
 - $SC(A, r) = \frac{n_r}{V(A, r)}$ if distinct values distributed evenly

Size of natural join

$r \times s$ contains $n_r n_s$ tuples and each tuple of $r \times s$ occupies $s_r + s_s$ bytes

1. If $R \cap S = \emptyset$, $r \bowtie s = r \times s$
2. If $R \cap S$ is a key for R , then $r \bowtie s = \text{number of tuples in } s$
3. If $R \cap S = \{A\}$ is not a key for neither, $r \bowtie s = \frac{n_r n_s}{\min\{V(A, r), V(A, s)\}}$

Selection Operation

- Linear Search = $b_r(\text{read block}) + \frac{b_r}{2}(\text{selection on key attribute})$
- Binary Search = $\lceil \log_2(b_r) \rceil (\text{locate first tuple}) + \lceil \frac{SC(A, r)}{f_r} - 1 \rceil (\text{number of satisfying records})$
- Clustered index, equality on key = index access + 1
- Clustered index, equality on nonkey = index access + $\lceil \frac{SC(A, r)}{f_r} \rceil$
- Unclustered index, equality on key = index access + $SC(A, r)$

Example

Assuming the following sizes

Sailors : 500 pages, 80 tuples/page, 50 bytes/tuple

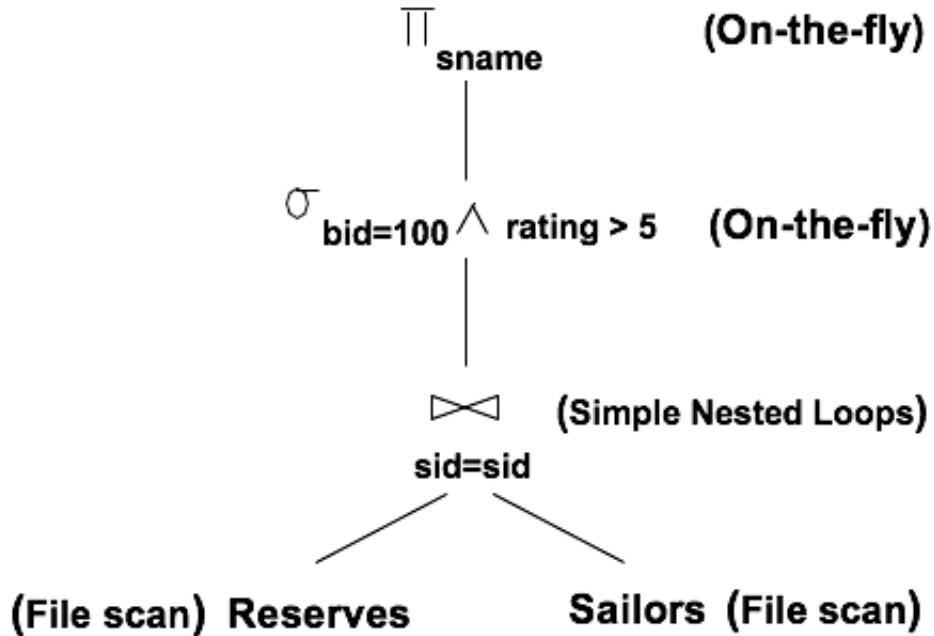
Reserves : 1000 pages, 100 tuples/page, 40 bytes/tuple

```

1  SELECT S.name
2  FROM Reserve R, Sailors S
3  WHERE R.sid=S.sid AND R.bid=100 AND S.rating>5

```

Relational Algebra : $\Pi_{sname}(\sigma_{bid=100 \wedge rating>5}(Reserves \bowtie_{sid=bid} Sailors))$



$$\text{Simple Nested Loops} = 500 + 500 \times 1000 = 500500 \text{ I/Os}$$

On-the-fly Evaluation

- The evaluation of Op_2 is **on-the-fly** if the result of Op_1 is **directly sent** to Op_2
NOT stored in a temporary file and without disk access
- Also called **pipelined evaluation**

Materialized Evaluation

- The result is stored in a temporary file first
- Required more disk accesses

Concurrency Control

- Transaction is atomic and a sequence of read/write operations

```

1 | t1: read(a)           t2: a = 100
2 |   a = a + 100         write(a)
3 |   write(a)            read(a)
  
```

$$T_1 = r_1[a]w_1[a] \quad T_2 = w_2[a]r_2[a]$$

Serial Schedule

- The operations belonged to same single transaction appear together

$$T_1 : r_1[a]w_1[a]$$

$$T_2 : \quad \quad \quad w_2[a]r_2[a]$$

Serializable Schedule

- The effect is equivalent to some serial schedules

$T_1 : r_1[a] \quad w_1[a]$

$T_2 : \quad w_2[b] \quad r_2[a]$

Equivalent to $T_1 T_2$

$T_1 : r_1[a] w_1[a]$

$T_2 : \quad w_2[b] r_2[a]$

View Serializability

- Two schedules H and H' are view equivalent if
 - Same set of transactions and same operations
 - For any x , if T_i reads x from T_j in H , then also in H'
 - For each x if $w_i[x]$ is final write of x in H then also in H'
- A schedule is view serializable if it is **view equivalent to a serial schedule**

$r_3[x] w_4[x] w_3[x] w_6[x]$

- T_3 read x from T_b .
- The final write for x is $w_6[x]$.
- View equivalent to $T_3 T_4 T_6$.

$r_3[x] w_4[x] w_3[x]$

- T_3 read x from T_b .
- The final write for x is $w_3[x]$.
- Not serializable.

$r_3[x] w_4[x] r_7[x] w_3[x] w_7[x]$

- T_3 read x from T_b .
- T_7 read x from T_4 .
- The final write for x is $w_7[x]$.
- View equivalent to $T_3 T_4 T_7$.

$w_1[x] r_2[x] w_2[x] r_1[x]$

- T_2 read x from T_1 .
- T_1 read x from T_2
- The final write for x is $w_2[x]$.
- Not serializable.

Conflict Serializability

- Non-conflicting operations can always be **swapped without affecting result** of other operations
- Swapable operations: $r_i[x] r_j[x] \iff r_j[x] r_i[x]$
- Two histories H and H' are **conflict equivalent** if
 - Same set of transactions and same operations
 - For any conflicting operations p_i and q_j , if p_i appears before q_j in H then same as H'
- A history is **conflict serialisable** if it is **conflict equivalent** to a serial history

Testing for conflict serializability = **Serialization Graph**

$V = \{T_1, \dots, T_n\}$ is set of transactions

E contains (t_i, t_j) if one of T_i operations precedes and conflicts with other

- It is conflict serializable iff there exists a DAG in serialization graph

Lock-based Protocols

- Two kinds of locks
 - Shared-locks for read operations
 - Exclusive-locks for write operations
- Shared-lock allowed read but not write operations
- Exclusive lock does not allow both read and write operations

Two-phase Locking Protocol

- Each transaction issues lock and unlock requests in two phases
 - Growing phase: A transaction **may obtain locks** but may **not release any lock**
Lock can be **upgrade** in the growing phase
 - Shrinking phase: A transaction **may release locks** but may **not obtain any new lock**
Lock can be **downgrade** in the shrinking phase

Recovery

Log-based Recovery

- Each record describes single database write and has following fields
 $< \text{Transaction name}, \text{item name}, \text{old value}, \text{new value} >$
- Some Special log records: $< T_i, \text{start} >$, $< T_i, \text{commits} >$, $< T_i, \text{abort} >$

Immediate Database Modification (undo/redo)

- Transaction T_i needs to **undo** if the log contains $< T_i, \text{start} >$ but not contain
 $< T_i, \text{commits} >$
- Transaction T_i needs to **redo** if the log contains both $< T_i, \text{start} >$ and $< T_i, \text{commits} >$

Deferred Database Modification (No-undo/redo)

- Each record describes single database write and has following fields
 $< \text{Transaction name}, \text{item name}, \text{new value} >$
- Transaction T_i needs to **redo** if the log contains both $< T_i, \text{start} >$ and $< T_i, \text{commits} >$